

Ανάπτυξη Λογισμικού για Δυσεπίλυτα Αλγοριθμικά Προβλήματα

Ενότητα 2: Clustering

Γιάννης Εμίρης

Τμήμα Πληροφορικής & Τηλεπικοινωνιών
Πανεπιστήμιο Αθηνών

November 10, 2016

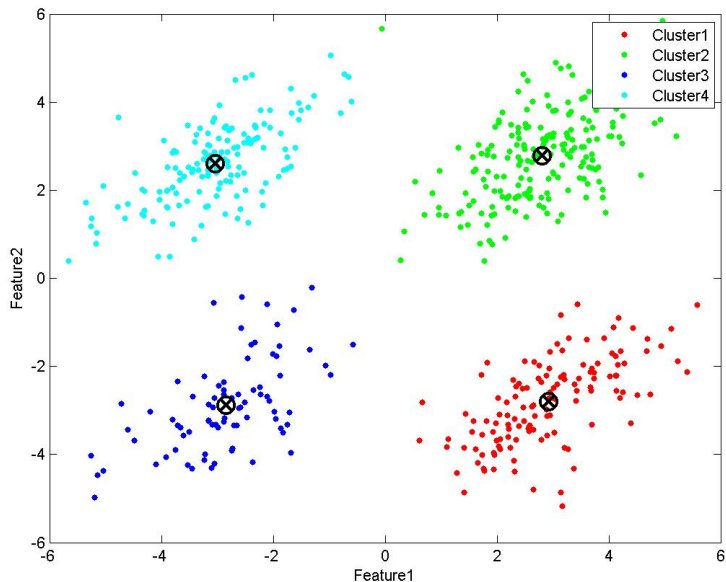
- 1 Clustering
- 2 Vector spaces
 - Elkan's Algorithm
- 3 Metric spaces
- 4 General Improvements
 - Swapping
 - Sampling
 - Initialization
 - Relation to LSH/DBH
- 5 Evaluation

Definition (k clusters)

Given n objects, and $k > 1$, partition the objects into k subsets (clusters) so as to optimize some objective function.

- Objects in the same cluster are more "similar" (or closer) to each other than to those in other clusters.
- Possible criteria: minimizing the total distance among all cluster points, minimizing the distance of cluster points to some center, etc.
- Variations: k is unknown and computed, e.g., by the Silhouette method. Capacitated/balanced: k given, clusters of equal cardinality.
- Applications: Classification, Social Network Analysis, Recommender Systems, Market Research, Bioinformatics etc

Good Clustering, with centers



- hierarchical (agglomerative): each point initializes a cluster, merge until stopping criterion, e.g., predetermined number of clusters, or if merging creates cluster with points too far apart.
- point-assignment: given some initial clusters, assign points to "best" cluster; might allow combining / splitting clusters, or unassign points. Example: k-means (our focus).

(Ullman et al: Mining Massive datasets)

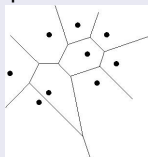
- 1 Clustering
- 2 Vector spaces
 - Elkan's Algorithm
- 3 Metric spaces
- 4 General Improvements
 - Swapping
 - Sampling
 - Initialization
 - Relation to LSH/DBH
- 5 Evaluation

Problem definition

- Clustering that minimizes objective function.
- k is given.
- Centroids do **not** have to be part of the dataset

k-means

- **k-means** is the most common problem: Main algorithms:
 - Lloyd's algorithm is standard.
 - Elkan's uses triangular inequality to accelerate updates.
- Also used to construct initial clusters for more sophisticated method.
- In Euclidean space, assignment is point location to the k Voronoi cells.



k-means: Objective function

Typical ambient space is \mathbb{R}^d but can generalize to metric space \mathcal{Z} .

Minimization function

In any metric space over points/objects \mathcal{Z} with distance/metric function d , let the dataset be $X = \{x_1, \dots, x_n\} \subseteq \mathcal{X} \subseteq \mathcal{Z}$, $k > 1$. Given centroids $C \subset \mathcal{Z}$, let

$$d(x_i, C) = \min_{c \in C} d(x_i, c).$$

Consider vector $v(C) = (d(x_1, C), \dots, d(x_n, C))$. The k -means objective is:

$$\min_{C \subseteq \mathcal{Z}, |C|=k} \|v(C)\|_2^2 = \sum_{i=1}^n d(x_i, C)^2.$$

The k -means objective is NP-hard, but for the ℓ_2 metric, Lloyd's algorithm converges quickly to a *local* minimum.

Various minimizations

Recall $X = \{x_i\}$, $v(C) = (d(x_1, C), \dots, d(x_n, C))$, where $C \subset \mathcal{Z}$ are the centroids, and the k -means objective is:

$$\min_{C \subseteq \mathcal{Z}, |C|=k} \|v(C)\|_2^2 = \sum_{i=1}^n d(x_i, C)^2.$$

Similar objectives:

- k -median: $\min_{C \subseteq \mathcal{Z}, |C|=k} \|v(C)\|_1$,
- k -medoid: $\min_{C \subseteq X, |C|=k} \|v(C)\|_1$.
- k -center: $\min_{C \subseteq X, |C|=k} \|v(C)\|_\infty$,

Algorithm

Initialize k centers randomly (or using some strategy).

- 1 Assignment: Assign each object to its nearest center.
- 2 Update: Calculate mean of each cluster, make it the new center.

Repeat the two steps until there is no change in the assignments.

Properties

- Each distance calculation = $O(d)$ because vectors in \mathbb{R}^d .
- Assignment = $O(nkd)$, Update = $O(nd)$,
- #iterations unknown, in practice $\ll n$.
- Converges to local minimum in Euclidean space (depends on initialization)

- 1 Clustering
- 2 Vector spaces
 - Elkan's Algorithm
- 3 Metric spaces
- 4 General Improvements
 - Swapping
 - Sampling
 - Initialization
 - Relation to LSH/DBH
- 5 Evaluation

Lemma

For centers a, b , and point x : $d(x, a) \leq d(a, b)/2 \Rightarrow d(x, a) \leq d(x, b)$.

Proof. $x \in \text{Ball}$ centered at a with radius $d(a, b)/2$.

Corollary: x shall not be assigned to b .

Lemma (Triangular)

If center c is updated to c' then:

$$|d(x, c) - d(c, c')| \leq d(x, c') \leq d(x, c) + d(c, c').$$

Initialisation

- Pick initial centers c and assign each x to closest center $c(x)$.
- Set $u(x) \leftarrow d(x, c(x))$, $l(x, c) \leftarrow d(x, c)$, $\forall c \neq c(x)$.
- Compute $d(c, c')$ for all centers $c \neq c'$.

Update bounds

Compute $d(c, c')$ for all centers $c \neq c'$.

- Upper bound $u(x)$ on distance of x to its center c ;
for new center c' of x 's cluster: update $u(x) \leftarrow u(x) + d(c, c')$.
- Lower bound $l(x, c)$ to any other center $c \neq c(x)$;
update when c changes to c' : $l(x, c') \leftarrow |l(x, c) - d(c, c')|$.

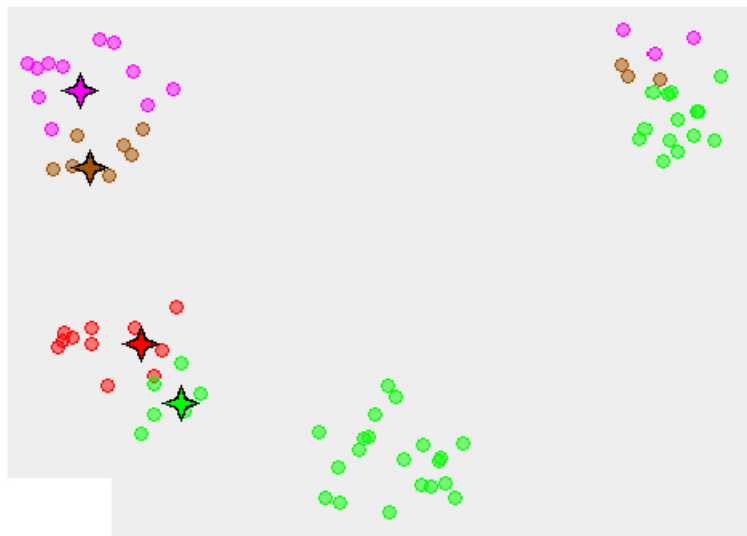
These bounds follow from the triangular inequalities.

```
for each point  $x$  do  
  if  $u(x) \leq d(c(x), c')/2$ , for all centers  $c' \neq c(x)$  then Skip  $x$ ; continue  
  end if  
  for each  $c'$  to which  $x$  is not assigned do  
    if  $u(x)$  not opt, compute  $u(x) \leftarrow d(x, c(x))$ ; mark  $u(x)$  as opt;  
    if  $u(x) \leq d(c, c')/2$ , or  $u(x) \leq l(x, c')$  then Skip  $c'$ ; continue  
    end if  
    compute  $d(x, c')$ ;  
    if  $d(x, c') < u(x)$  then assign  $x$  to  $c'$ ; set  $u(x) \leftarrow d(x, c')$   
    end if  
    update  $l(x, c') \leftarrow d(x, c')$ ;  
  end for  
end for
```

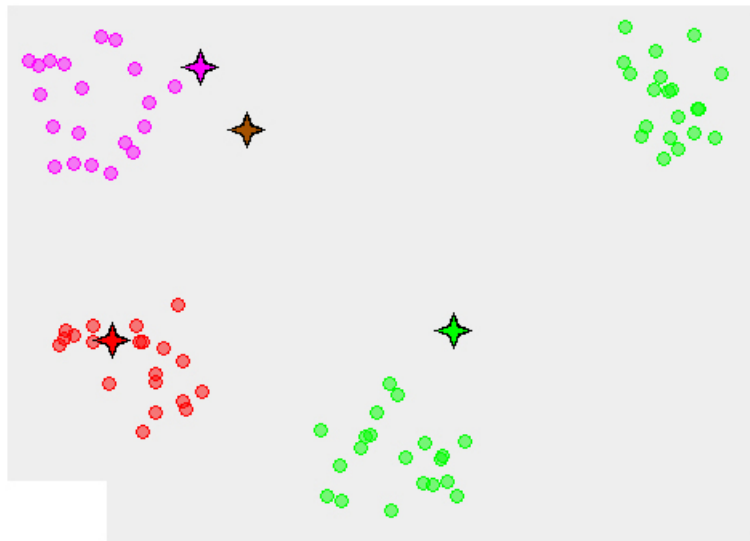
Elkan's algorithm offers, compared to Lloyd's,


- large experimental speedup,
- same output, same convergence,
- requires much higher storage.


Initialization



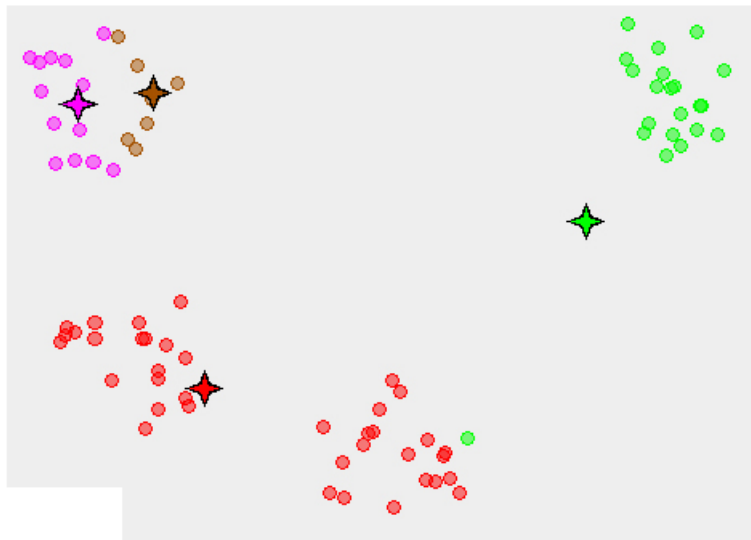
Lloyd's 1

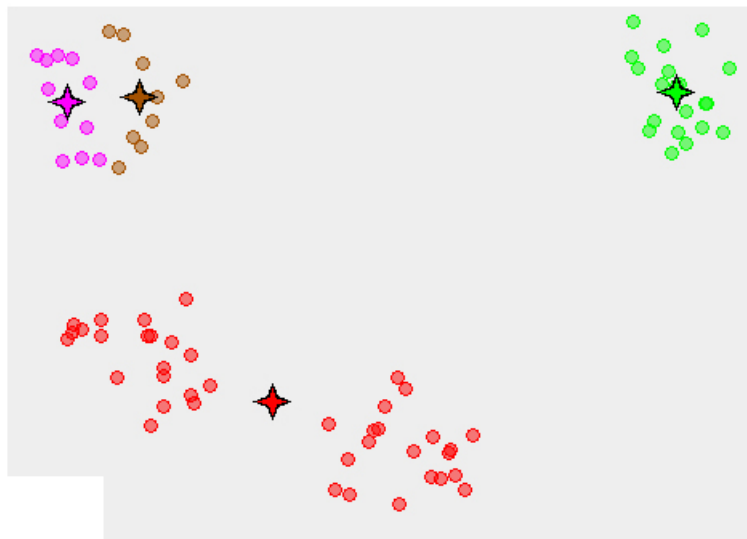






Lloyd's 4





- 1 Clustering
- 2 Vector spaces
 - Elkan's Algorithm
- 3 Metric spaces
- 4 General Improvements
 - Swapping
 - Sampling
 - Initialization
 - Relation to LSH/DBH
- 5 Evaluation

Goal: Handle any distance metric; k-means only if consistent with mean.

k-medoids (PAM is simplest algorithm) use centroids that **belong** to the dataset:

Definition (Medoid)

The medoid of a set is the object of the set that minimizes total dissimilarity (distance) to all other objects in the set.

Objective function (cf. above): Minimize sum of distances to point's centroid.

vs k-means

k-means tends to select convex spherical clusters; k-medoids less so.

k-means is more sensitive to noisy data and outliers.

k-means is faster and easier to implement.

(Kaufman-Rousseeuw'87)

Partitioning Around Medoids (PAM)

Initialize k centroids randomly.

① Assignment: Assign each object to nearest centroid; compute objective

② Update:

for each centroid m **do**

for each non-centroid t **do**

 Swap m and t , compute new value of objective function.

 Update the configuration if new objective value is lower.

end for

end for

Update until no new configuration selected.

Let distance calculation = $O(d')$. Update of Objective = $O((n - k)d')$ if 2nd best centroid known. Hence, update = $O((n - k)^2 kd') \sim n^2$.

(Small improve: test candidate swaps, implement best one.)

Update of Objective cost function

$$\text{Objective function: } J = \sum_{i=1}^{n-k} \text{dist}(i, c(i)).$$

Suppose for each item i we also store the 2nd best centroid c' .


If centroid m is replaced by non-centroid t , ΔJ takes 4 possible values:

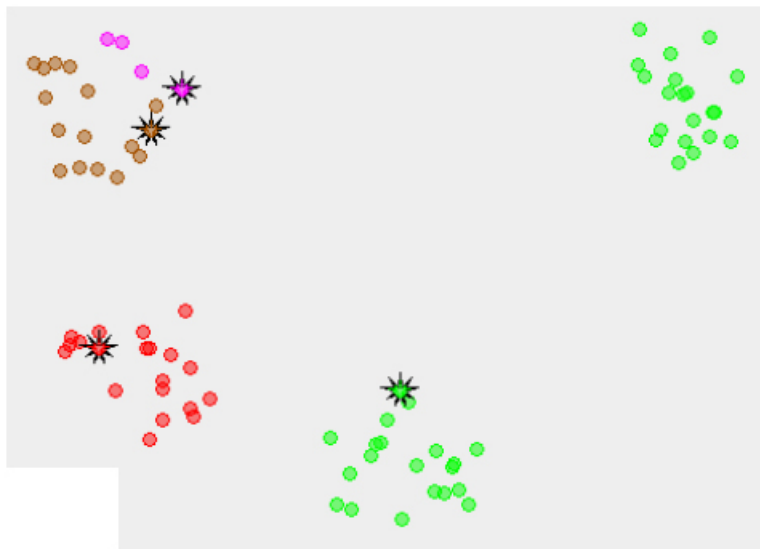
- For $i : c(i) = m$,

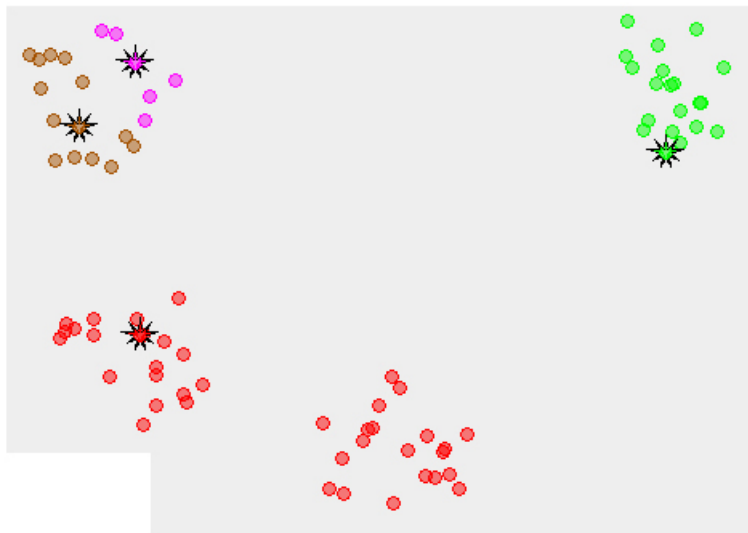
$$\Delta J = \begin{cases} \text{dist}(i, t) - \text{dist}(i, m), & \text{if } \text{dist}(i, t) \leq \text{dist}(i, c') : \text{do nothing} \\ \text{dist}(i, c') - \text{dist}(i, m), & \text{if } \text{dist}(i, t) > \text{dist}(i, c') : \text{assign } i \text{ to } c' \end{cases}$$


- For $i : c(i) \neq m$,

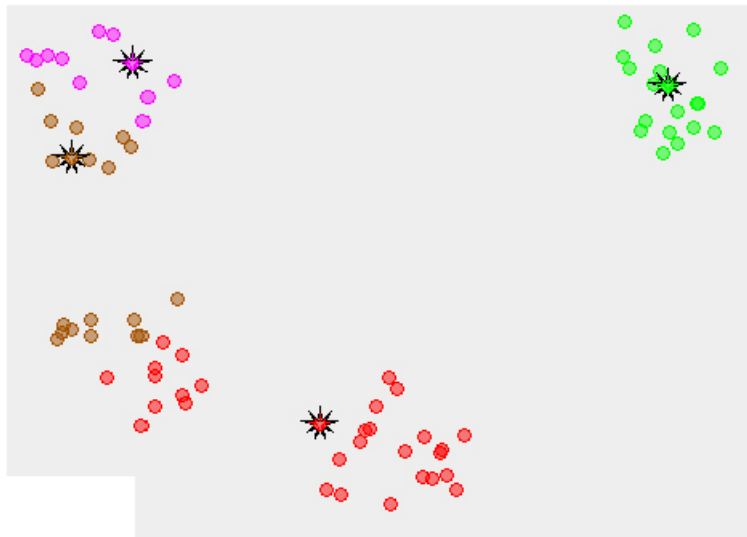
$$\Delta J = \begin{cases} 0, & \text{if } \text{dist}(i, t) \geq \text{dist}(i, c(i)) : \text{do nothing} \\ \text{dist}(i, t) - \text{dist}(i, c(i)), & \text{if } \text{dist}(i, t) < \text{dist}(i, c(i)) : \text{assign } i \text{ to } t \end{cases}$$




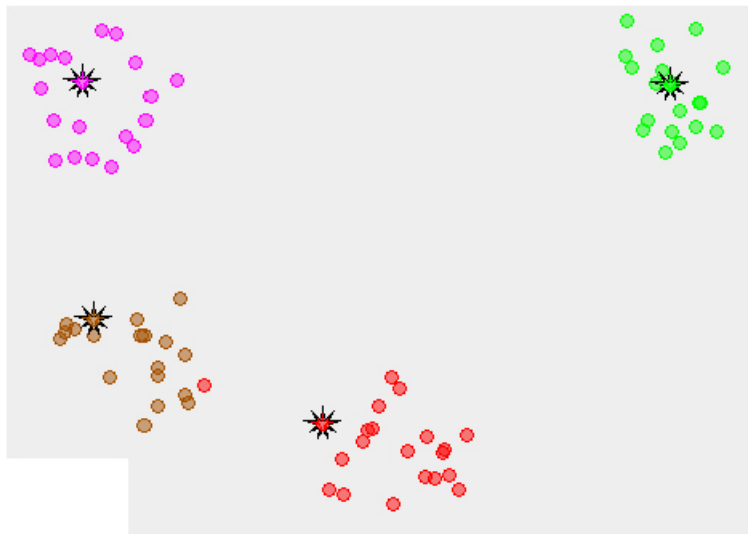


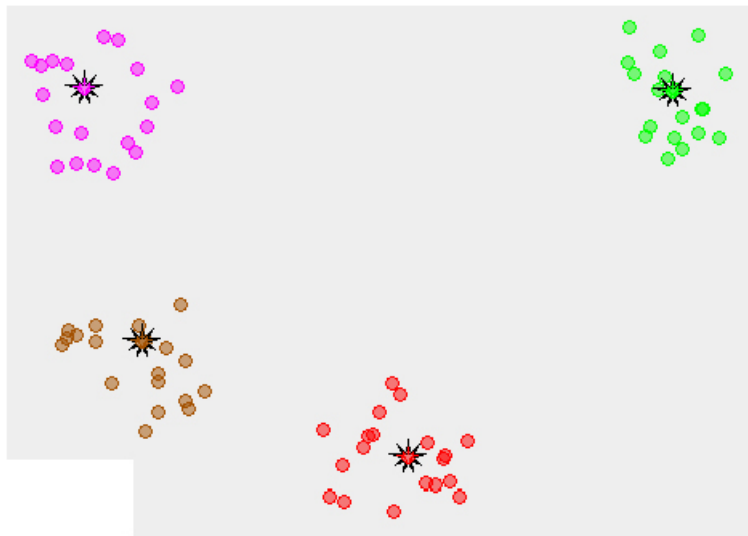












- 1 Clustering
- 2 Vector spaces
 - Elkan's Algorithm
- 3 Metric spaces
- 4 General Improvements**
 - Swapping
 - Sampling
 - Initialization
 - Relation to LSH/DBH
- 5 Evaluation

- 1 Clustering
- 2 Vector spaces
 - Elkan's Algorithm
- 3 Metric spaces
- 4 General Improvements**
 - **Swapping**
 - Sampling
 - Initialization
 - Relation to LSH/DBH
- 5 Evaluation

Accelerating updates

Two faster updates, which may however lose accuracy compared to PAM.
Recall that after every swap we compute J in $O((n - k)d')$.

1. Improved Update

Instead of swapping centroid m with every point t , swap m only with every non-centroid in same cluster as m .

Complexity: $n - k$ iterations instead of $k(n - k)$, hence update = $O((n - k)^2 d')$

2. Update à la Lloyd's

For every cluster: (i) compute its medoid t , (ii) Swap its current centroid m with t .

The medoid t minimizes $\sum_{i \in C} d(i, t)$ over all possible objects t in cluster C .
Computed in $O(a^2 d')$, assuming clusters have $a \simeq n/k$ items.

Total Complexity = $O((ka^2 + k(n - k))d') = O((n^2/k + nk)d') = O(n^2 d')$
(Park-Jun'09).

- 1 Clustering
- 2 Vector spaces
 - Elkan's Algorithm
- 3 Metric spaces
- 4 General Improvements**
 - Swapping
 - Sampling**
 - Initialization
 - Relation to LSH/DBH
- 5 Evaluation

Clustering Large Applications (CLARA)

Idea: work with sample of size $n' < n$. Use s samples drawn independently.

Overall algorithm:

for $i = 1, \dots, s$ **do**

 apply PAM on a random (uniform) sample of size n'

 assign n points to k computed centroids

 calculate the total cost of the partitioning

end for

return best partitioning

Experimental results recommend: $s = 5, n' = 40 + 2k$.

CLARA based on RANdomised Search (CLARANS)

- Update: swap m 's with t 's, for some randomly selected (m, t) 's only.
- Implement by picking random $Q \subset \{1, \dots, k\} \times \{1, \dots, n - k\}$.
- Also try s different, independent initializations.

for $i = 1, \dots, s$ **do**

Initialize k centroids (independently of previous initial sets).

Randomly select set Q of $|Q|$ pairs (m, t) , $|Q| \ll n(n - k)$.

Run PAM swapping centroid m with non-centroid t , $\forall (m, t) \in Q$.

Keep configuration with minimum objective value.

end for

Output clustering with minimum objective value over s candidates.

Experimental results recommend: $s = 2$, $|Q| = \max\{0.12 \cdot k(n - k), 250\}$.
(Ng-Han:VLDB'94) (Theodoridis-Koutroumbas:Pattern Recognition,ch.14)

Implement CLARANS

Consider a dataset of n points/objects with k clusters.

How to choose one uniformly distributed pair from

$$\{(m, t) : m = 0, \dots, k - 1, t = 0, \dots, n - 1\}.$$

Pick a uniformly distributed integer $x \in [0, kn - 1]$ and return

$$(x \bmod k, \left\lfloor \frac{x}{k} \right\rfloor).$$

CLARA / CLARANS are approximations of PAM.

CLARANS is in $O(n^2)$, CLARA is significantly faster.

CLARANS leads to a better clustering, especially if CLARA's random sample is biased.

Both methods rely heavily on their parameters.

- 1 Clustering
- 2 Vector spaces
 - Elkan's Algorithm
- 3 Metric spaces
- 4 General Improvements**
 - Swapping
 - Sampling
 - Initialization**
 - Relation to LSH/DBH
- 5 Evaluation

Improve Initialisation 1: Spread-out

initialization++ : k-means++ / k-medoids++:

- (1) Choose a centroid uniformly at random; $k \leftarrow 1$.
- (2) \forall point i , $D(i) \leftarrow$ min distance to centroid among k chosen centroids
- (3) Choose new centroid: r chosen with probability proportional to $D(r)^2$:

$$\text{prob}[\text{choose } r] = D(r)^2 / \sum_{j=1}^{n-k} D(j)^2.$$

Let $k \leftarrow k + 1$.

- (4) Go to (2) until $k =$ given #centroids.

Expected approximation ratio = $O(\log k)$ (Arthur-Vassilvitskii:SODA'07)

Implement initialization++

Given $D(i) > 0$, $i = 1, \dots, n$, compute $n - k$ partial sums

$$P(r) = \sum_{j=1}^r D(j)^2, \quad r = 1, \dots, n - k,$$

and store them in an array P of $n - k$ `float` numbers. To avoid the $P(r)$'s being very large, one can normalize all $D(i)$'s, e.g. by dividing them by $\max_i D(i)$.

Pick a uniformly distributed `float` $x \in [0, P(n)]$ and return

$$r \in \{1, 2, \dots, n - k\} : P(r - 1) < x \leq P(r),$$

where $P(0) = 0$: r chosen with probability proportional to $P(r) - P(r - 1) \sim D(r)^2$. Can find r by binary search in array P .

Improve Initialisation 2: Concentrate

Select centroids close to dataset's center of mass (and to each other) as follows.

(1) Calculate symmetric $n \times n$ distance matrix of all objects, i.e. all distances d_{ij} from every object $i = 1, \dots, n$ to every other object $j = 1, \dots, n, i \neq j$.

(2) For object i compute

$$v_i = \sum_{j=1}^n \frac{d_{ij}}{\sum_{t=1}^n d_{jt}}, \quad i = 1, \dots, n.$$

(3) Return the k objects with k smallest v_i values.

Algorithm proposed in (Park-Jun'09).

- 1 Clustering
- 2 Vector spaces
 - Elkan's Algorithm
- 3 Metric spaces
- 4 General Improvements**
 - Swapping
 - Sampling
 - Initialization
 - **Relation to LSH/DBH**
- 5 Evaluation

Direct approach

At each iteration:

- 1 Index k centroids into data-structure (e.g. LSH hashtables)
- 2 For every point, run ANN to find nearest centroid.

Reverse approach

- Index n points into L hashtables: once for entire algorithm.
- LSH/DBH TableSize $\leq n/8$: avoid buckets with very few items.
- At each iteration, for each centroid c , range/ball queries centered at c .
- Mark assigned points: either move them at end of LSH buckets (and insert "barrier"), or just flag them (using a new "flag" field).
- Increase radii by $\times 2$, start with $\min(\text{dist between centers})/2$, until all points assigned, or most ranges/balls do not assign a new point.
- If point in ≥ 2 ranges/balls compare its distances to respective centroids
- At end: for every unassigned point, compare its distances to all centroids

- 1 Clustering
- 2 Vector spaces
 - Elkan's Algorithm
- 3 Metric spaces
- 4 General Improvements
 - Swapping
 - Sampling
 - Initialization
 - Relation to LSH/DBH
- 5 Evaluation

Evaluate clustering without reference to objective function. Try to capture meaning of clustering.

- Let k be the number of computed clusters.
- Internal evaluation considers the given pointset and the clusters, produces quality coefficient for each partition; k may be a parameter.
- External evaluation: use known class labels and benchmarks; often created by humans.

In the sequel we present internal evaluation methods, mainly Silhouette.

For $1 \leq i \leq n$, $a(i)$ = average distance of i to other objects in same cluster.
Let $b(i)$ = average distance of i to objects in *next best* (neighbor) cluster.

Silhouette of Object i

$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}} = \begin{cases} 1 - a(i)/b(i), & \text{if } a(i) < b(i) \\ 0, & \text{if } a(i) = b(i) \\ b(i)/a(i) - 1, & \text{if } a(i) > b(i) \end{cases} \in [-1, 1].$$

Interpret silhouette

- $s(i) \rightarrow 1$: i seems correctly assigned to its cluster;
- $s(i) \simeq 0$: borderline assignment (but not worth to change);
- $s(i) \rightarrow -1$: i would be better if assigned to next best cluster.

Silhouette: Cluster and clustering

Specific clusters

Evaluate a cluster: Compute average $s(i)$ over all i in some cluster.

If k is too large or too small, some clusters shall display much smaller silhouettes than the rest. Silhouette plots are used to improve k : try different k 's and see if clusters have roughly equal silhouettes.

Overall Clustering

Overall Silhouette coefficient = average $s(i)$ over $i = 1, \dots, n$.

High if well clustered, low may indicate bad k (or existence of outlier points).